

Towards Zero Unknown Word in Neural Machine Translation

Xiaoqing Li,[†] Jiajun Zhang,[†] Chengqing Zong^{†‡}

[†] National Laboratory of Pattern Recognition, Institute of Automation,
Chinese Academy of Sciences

[‡] CAS Center for Excellence in Brain Science and Intelligence Technology
{xqli,jjzhang,cqzong}@nlpr.ia.ac.cn

Abstract

Neural Machine translation has shown promising results in recent years. In order to control the computational complexity, NMT has to employ a small vocabulary, and massive rare words outside the vocabulary are all replaced with a single *unk* symbol. Besides the inability to translate rare words, this kind of simple approach leads to much increased ambiguity of the sentences since meaningless *unks* break the structure of sentences, and thus hurts the translation and reordering of the in-vocabulary words. To tackle this problem, we propose a novel substitution-translation-restoration method. In substitution step, the rare words in a testing sentence are replaced with similar in-vocabulary words based on a similarity model learnt from monolingual data. In translation and restoration steps, the sentence will be translated with a model trained on new bilingual data with rare words replaced, and finally the translations of the replaced words will be substituted by that of original ones. Experiments on Chinese-to-English translation demonstrate that our proposed method can achieve more than 4 BLEU points over the attention-based NMT. When compared to the recently proposed method handling rare words in NMT, our method can also obtain an improvement by nearly 3 BLEU points.

1 Introduction

Neural machine translation is a recently proposed approach to MT and has shown competing results to conventional translation methods [Kalchbrenner and Blunsom, 2013; Cho *et al.*, 2014; Sutskever *et al.*, 2014]. In neural machine translation, the source sentence is converted into vector representation by a neural network called encoder, then another neural network called decoder generate target sentence word by word based on source representation and target history. This framework has several advantages over conventional translation methods. First, it does not need any domain knowledge as required by conventional methods to design features. Second, the distributed representation allows NMT model to generalize well and produce novel translations for source words and phrases, while the symbolic representation in conventional MT makes

it impossible to generate translations beyond the rule table extracted from the bilingual corpus. Third, the memory consumption of NMT model is also much smaller.

Despite these advantages, NMT models have a major drawback in handling rare words. In order to control the computational complexity, which grows proportional to target vocabulary size¹, most NMT systems limit the vocabulary to contain only 30k to 80k most frequent words in both the source and target side and convert rare words into a single *unk* symbol. An obvious problem of this approach is that NMT model cannot learn the translation of rare words. In particular, if a source word is outside the source vocabulary or its translation is outside the target vocabulary, the model will not be able to generate proper translation for this word during testing. Another problem is that masking rare words with meaningless *unk* will increase the ambiguity of the sentence. This can be illustrated by the following three sentences,

- a) Mike chases the pet with *mottle*
- b) Mike chases the pet with *scooter*
- c) Mike chases the pet with *Sullivan*

Assume all the last words in the three sentences are rare words. The word 'mottle' in sentence a) modifies the object 'pet', and both the word 'scooter' and 'Sullivan' in sentence b) and c) modifies the predicate, but one describes the tool and the other describes the companion. The translation of the preposition 'with' and the word order will be quite different when translating the three sentences into Chinese. If the last words are replaced by the *unk* symbol, the three sentences will be the same. As a result, the model can only generate the translation by chance.

To solve the above problems, we propose a novel rare word replacement method based on similarity. During training, word alignment will first be induced from bilingual corpus. And each aligned word pair which contains rare word either on the source side or the target side will be replaced with similar in-vocabulary words, where the similarity model is learned from a large mono-lingual corpus. Then this new bilingual corpus with rare words replaced will be used to train a NMT model. During testing, the rare words in input sentence will also be replaced with similar in-vocabulary words.

¹source vocabulary size contributes less to computational complexity, but knowing how to translate source word to target *unk* is not helpful, so the source vocabulary size is also limited.

After translation, a post-processing step is adopted to recover the translation of rare words.

Experiments on Chinese to English translation task show that more than 4 points in BLEU score can be gained with our approach over the baseline. And the gain is also much larger than a previously proposed replacement method [Luong *et al.*, 2015b].

2 Neural Machine Translation and Impact of Rare Words

In this section, we first give a brief introduction to neural machine translation and explain why NMT model could not employ large vocabulary. Then we quantitatively analyze how rare words impact the performance of NMT.

2.1 Neural Machine Translation

Neural machine translation is conceptually simple: it models the translation probability of a source sentence $s = (s_1, s_2, \dots, s_m)$ into target sentence $t = (t_1, t_2, \dots, t_n)$ with a single neural network as follows,

$$p(t|s) = \prod_{i=1}^n p(t_i|t_{<i}, s)$$

where the conditional probability is often parameterized with the encoder-decoder framework. The encoder reads the source sentence and encodes it into a sequence of hidden states $h = (h_1, h_2, \dots, h_m)$:

$$h_i = f(s_i, h_{i-1})$$

Then the decoder generates the translation word by word based on the target hidden states $z = (z_1, z_2, \dots, z_n)$:

$$p(t_i|t_{<i}, s) = \frac{1}{Z} \exp \{q(t_i, t_{i-1}, z_i, c_i)\}$$

where

$$z_i = g(t_{i-1}, z_{i-1}, c_i)$$

$$c_i = r(z_{i-1}, h)$$

In above formulations, f, g, g and r are non-linear transformations and varies in different systems.

The most time consuming step in the network is the calculation of the normalization constant Z , which is computed as follows,

$$Z = \sum_{t' \in IV} \exp \{q(t', t_{i-1}, z_i, c_i)\}$$

According to this equation, we need to iterate over all target in-vocabulary words to calculate a non-linear transformation for each, and then sum them up². So the total computational complexity will grow almost proportional to the target vocabulary size. Considering that it usually takes days to weeks to

²there is no normalization over the input vocabulary, and the operation corresponding to each source word is indexing rather than non-linear transformation, so the source vocabulary size has much smaller influence on computational complexity




Figure 1: Impact of missing translations and increased ambiguity
red line: performances (BLEU) of sentence groups with different number of rare words
blue line: performances after setting different number of words to *unk* in translations of the sentence group without rare words

train a NMT model on a large corpus with a vocabulary size of 30k to 80k, training with the whole target vocabulary is obviously infeasible. So addressing the problem for rare words is quite necessary for neural machine translation.

2.2 Impact of Rare Words

As discussed in the introduction part, rare words cause two problems for neural machine translation. First, NMT model cannot learn translations for rare words because they are all converted to *unk* in the training data. Second, rare words increase the ambiguity of the sentence, which increases the difficulty to translate and reorder the rest in-vocabulary words in the sentence.

To quantitatively check the impact of the two factors, we design the following experiment. We extract 5 groups of Chinese sentences with different number of rare words (0-4) from the NIST Chinese to English translation data set. Each group contains 50 sentences together with their reference translations. In order to rule out the influence of sentence length, all the sentences in the 5 groups are between 20 to 30 words. We use the same system to translate these sentences and the corresponding performances are shown in Figure 1 (red line). It is obvious more rare words lead to worse performance. To simulate the impact of missing translation for rare words, we randomly set 1-4 words to *unk* in the translation of sentences in group 0. The result is shown as blue line. It could be inferred that the remaining gap between the red line and the horizontal line (denoting the performance of group 0) is caused by the increased ambiguity. According to the figure, when there are only one rare word in the sentence, the performance drop is mainly caused by missing translations, but when there are more rare words, increased ambiguity also contributes a lot to the performance drop.

3 Replace Rare Words with Similar words

The analysis in the above section shows the importance of keeping the sentence structure complete. So we propose

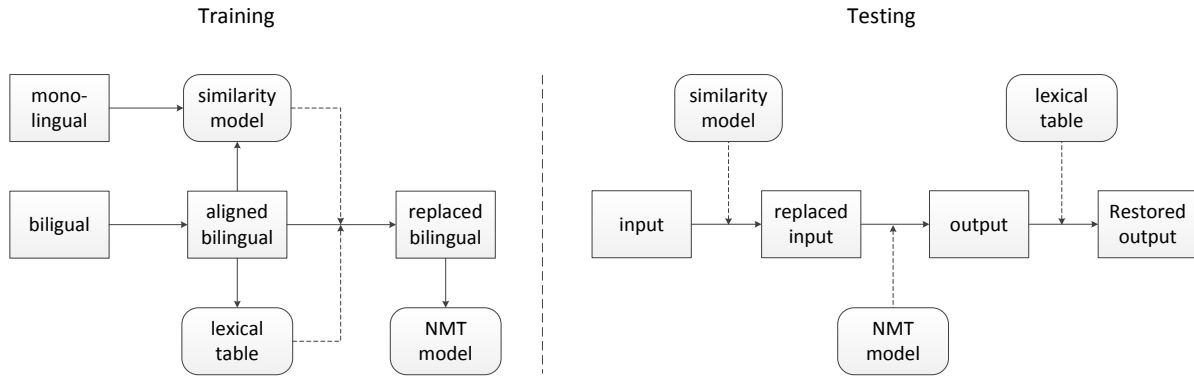


Figure 2: Data processing diagram for training and testing

to replace rare words in training and testing data with in-vocabulary words similar to them. The data processing diagram is shown in Figure 2.

In the training phase, we first learn a similarity model from a monolingual corpus, which is used to evaluate the similarity between words. We also need to learn word level alignment for sentence pairs in the bilingual corpus. As a byproduct, a lexical translation table can be derived from the aligned bilingual corpus. In our experiments, we only reserve the translation with the highest probability for each word in the table. Then the aligned word pairs which contain rare words will be replaced with in-vocabulary words similar to them. Finally, a NMT model will be learned from the new bilingual corpus.

In the testing phase, the rare words in testing sentence will be first replaced with similar in-vocabulary words. Then the sentence after replacement will be translated by the NMT model obtained in the training phase. With the help of the lexical translation model, the translation of those rare words will be substituted back into the generated target sentence to obtain the final result.

There are three issues not explained in detail in the diagram, including i) which words in the bilingual corpus will be replaced? ii) how to evaluate similarity between words? iii) How to recover the translation for rare words during testing? The following parts in this section will answer these questions.

3.1 Words to Be Replaced

Different languages are not perfectly corresponded in word level. For example, English articles are usually omitted when translated into Chinese. And the city name New York is just one word in Chinese. Sometimes the correspondence is even at phrase or sentence level, such as the translation of idioms. In this paper, we only handle word pairs with one-to-one mapping and rare words aligned to null. According to whether the source and the target word is rare, there are five cases.

- *unk* to *unk*, both the source and target word in the aligned pair are rare words. In this case we will replace the source word with a similar in-vocabulary word and the target word with the translation of the similar word.
- *unk* to common, only the source word is rare. In this case we will keep the target word and replace the source

word with the translation of the target word.

- common to *unk*, only the target word is rare. In this case we will keep the source word and replace the target word with the translation of the source word.
- common to common, no replacement in this case.
- *unk* to null or null to *unk*, source or target rare word is not aligned to any word. In this case we simply remove the rare word from the sentence.

3.2 Similarity Model

Distributed word representation has been shown powerful to capture syntactic and semantic information about words, and it is widely applied in various tasks [Turian *et al.*, 2010]. We adopt it here to find the most similar word for a given word w as follows,

$$w^* = \arg \max_{w' \in IV} sim(w, w')$$

in which IV denotes the set of in-vocabulary words, and the function sim is the cosine similarity between two word vectors.

$$sim(w, w') = \cos(vec(w), vec(w'))$$

However, since the word vectors and the lexical translation table are learned automatically from data, they may lead to inappropriate alternative for original translation pairs. For example, the most similar word to the rare word '善款' (donation) at the end of the following sentence is '筹募' (raise), which is in fact a synonym to the second to last word '募集'. As a result, this sentence will be ungrammatical after replacement because it will have two neighbouring predicates with the same meaning.

中国红十字会为新疆灾区募集善款
China Red-Cross for Xinjiang disaster-area raise donation

As another example, the similarity model find a synonym word '不和' to the rare word '失和' (discord) in the following sentence, but the lexical translation table gives it a wrong translation 'divorce'.

奥委会曾出现激烈的内部失和问题

IOC once appeared severe inner discord problem

To alleviate this problem, we propose to use multiple candidates provided by the similarity model, and choose one from them by checking whether it is fit for the bilingual context. Bi-directional ngram language model is adopted here for this purpose. For an aligned word pair (c_i, e_j) , the score to replace them with alternative (c'_i, e'_j) is calculated as follows,

$$\text{score} = p(c'_i|c_{i-1}, c_{i-2}) + p(c'_i|c_{i+1}, c_{i+2}) + \\ p(e'_j|e_{j-1}, e_{j-2}) + p(e'_j|e_{j+1}, e_{j+2})$$

The method to find the most appropriate alternative pair is described as follows. First, top N most similar words will be found for the source rare word. Then each of the source alternative together with its translation will be added to the candidate list. Finally, the bi-directional language model is used to rank these candidates and the best is adopted to replace the original translation pair.

As an alternative method to rerank the candidate pairs, we can also jointly consider bilingual word similarity. The similarity between the two translation pairs (c'_i, e'_j) and (c_i, e_j) will be calculated as follows,

$$\text{score} = \frac{\cos(\text{vec}(c'_i), \text{vec}(c_i)) + \cos(\text{vec}(e'_j), \text{vec}(e_j))}{2}$$

According to this measurement, only the translation pair which is similar to the original pair in both source and target side will be selected.

3.3 Restore Translation for Rare Words

Unlike traditional machine translation, in which output words and input words are explicitly linked by translation rules, the input and output in NMT are mapped in sequence level. Fortunately, the attention mechanism [Bahdanau *et al.*, 2015] provides a kind of soft alignment which can be used to find the corresponding source word for each target word. However, the alignment generated by current attention mechanism is far from perfect. Some source words are repeatedly attended and others are never attended. In order to reduce the influence of the alignment error, we add a constraint based on lexical translation table as follows. When a target word e_j aligns to a replaced source word c_i , and this pair can be found in the translation table, we will replace e_j with the translation of the original source word. Otherwise the target word e_j will be kept in the output.

4 Experiments

We evaluate our method on the Chinese to English translation task. Translation quality is measured by the BLEU metric [Papineni *et al.*, 2002].

4.1 Settings

The bilingual data to train the NMT model is selected from LDC, which contains about 0.6M sentence pairs. To avoid

spending too much training time on long sentences, all sentences pairs longer than 50 words either on the source side or on the target side are discarded. The alignment information needed for replacement are obtained by the Berkeley Aligner [Liang *et al.*, 2006] on the same bilingual data. We use the word2vec toolkit [Mikolov *et al.*, 2013] to train word vectors on the monolingual data, which is the combination of the source side of the bilingual data and Chinese Gigaword Xinhua portion. The Chinese bi-directional language model is trained with kenlm [Heafield *et al.*, 2013] on the same monolingual data, while the English language model is trained on the combination of the target side of the bilingual data and the English Gigaword.

The NIST 03 dataset is chosen as the development set, which is used to monitoring the training process and decide the early stop condition. And the NIST 04 to 06 are used as our testing set.

4.2 Training Details

The hyperparameters used in our network are described as follows. We limit both the source and target vocabulary to 30k in our experiments. This number of hidden units is 1,000 for both the encoder and decoder. And the word embedding dimension is 500 for all source and target words. The parameters in the network are updated with the adadelta algorithm.

To train the word vectors on monolingual data, we set the embedding dimension to 100 and the window size to 5. And we use top 10 most similar words in the similarity model considering bilingual context in section 3.2.

4.3 Main Results

We compare our best system (the one with bilingual similarity) to the baseline without any replacement, and the system proposed in [Luong *et al.*, 2015b], which only annotate target *unk* as *unk-k*, in which *k* indicates which source word translates into current unknown word. In particular, if e_j in the target sentence is a rare word and it's aligned to source word c_i , then *k* will be *i-j*. The performance of Moses [Koehn *et al.*, 2007] with 4-gram language model trained on the target side of the bilingual data is also shown for reference.

The results in Table 1 shows that our method significantly outperform the baseline by 4.15 BLEU points on average. It also surpasses the system proposed in [Luong *et al.*, 2015b] by 2.85 BLEU points. It's also worth to mention that the improvement given by their method is lower than the reported one on the French to English translation task. A possible reason is that there are much more reorderings in Chinese English language pair, so it's much harder to correctly predict which source word generate current target unknown word during translation. On the contrast, our model replaces rare words with similar words and keeps the completeness of the sentence, so that it is much easier for the translation model to learn correspondence between source and target words.

4.4 Comparison of Different Replacement Strategies

The performances of different replacement strategies are shown in Figure 3. It can be seen that considering bilingual context or bilingual similarity does improve the performance

System	03 (dev)	04	05	06	Average
Moses	28.68	29.87	27.27	29.17	28.77
Bahdanau et al. (2015)	25.65	28.94	25.13	27.86	26.90
Luong et al. (2015)	27.63	30.02	26.42	28.72	28.20
Ours	29.85	33.08	28.95	32.31	31.05

Table 1: Translation results for different systems. Bahdanau et al. (2015) is the NMT model with attention mechanism, which is adopted as our baseline without replacement. Luong et al. (2015) is the approach to decorate target *unk* with alignment information.




Figure 3: Comparison of different replacement strategies. The performance is an averaged one over all NIST data sets we adopt.

no-rep: use original sentence without replacement

simple: use most similar word for replacement

bi-lm: use bilingual bi-directional language model to choose from top similar words

bi-sim: use bilingual word similarity to choose from top similar words

over the simple replacement method, although the magnitude is not so significant when compared with the improvement of the simple method over the baseline.

We also show the performance of translating original testing sentences without replacement in the figure. The result is quite impressive. Nearly 3 BLEU points can be achieved over the baseline if we use the NMT model trained on the bilingual data with rare words replaced, while keeping the testing sentence unchanged. The improvement is even larger than that brought by replacing the testing data. This demonstrates that **training on complete sentences can greatly improve the quality of parameter estimation**, and thus lead to much better translations.

4.5 Better Attention after Replacement

It is also interesting to check how replacement affect the translation process in detail. Figure 4 shows the translations for the same sentence by different systems. The figure on the left corresponds to the baseline model without replacement. Because the third word '吁' (call) in the source sentence is outside the vocabulary, the baseline model cannot generate proper translation for it. What's more, wrong attention to this rare word results in bad translation for the common word

'美国' (America). The baseline model add an extra word 'north' before 'america', which is not a translation of any word in the source sentence. And the last source word '对话' (dialogue) is hardly attended by any target word, leading to missing translation for this word, although it is also a common word. On the contrary, our system find a similar word '呼吁' to the source rare word, which is in fact a synonym to it. Given this complete sentence without any rare word, our system is able to generate a nearly perfect translation for the source sentence, in which all source words are properly attended and translated. Since the rare word '吁' is not seen in the bilingual training corpus, the lexical translation table does not contain the translation for this word. So we keep the translation of the alternative word in the output. Last, the approach of [Luong *et al.*, 2015b] generated a similar translation as the baseline system (not shown in the figure). And even if a target *unk* is generated and aligned correctly to the source rare word, the translation of the rare word still cannot be restored because it's not in the lexical translation table.

4.6 Parameter Initialization

It is well known that parameter initialization has a big impact on the performance of neural networks. In this paper, we tried two ways to initialize the parameters of the system on replaced data. One is randomized initialization, the other is initializing with the parameter learned by the baseline model. According to our experience, the latter is robust and performs better than the former for our method. But for the approach proposed in [Luong *et al.*, 2015b], the latter initialization strategy does not bring any benefit.

5 Analysis of Untackled Rare Words

Although our method can handle more than 90 percent rare words in the data, there are still some remain untackled, which can be divided into two categories as follows.

One is related with complex alignments. As described in section 3.1, we only handle one to one and one to zero (zero to one) mapping in this paper. there are also some one to many (many to one) and many to many alignments in the data. Here is an example,

还要 标明 引进 批准 文号

and indicating the import ratification number

the rare word '文号' (document number) at the end of the source sentence aligns to two target word 'ratification' and 'number', and the target word 'ratification' also aligns to the second to last word '批准' (ratification). If we focus




Figure 4: Better attention after replacement. Darker block denotes larger attention weight. Left: translation by baseline model; Right: translation by our model.

on word level replacement, then replacing both 'ratification' and 'number' with the translation of a word similar to '文号' will make the source word '批准' unaligned. So it's better to do the replacement at phrase level. But how to find alternatives for phrases remains a problem and it will be leaved as our future work.

The other class of untackled rare words are related with the similarity model. According to Zipf's law [Zipf, 1949], it's impossible to contain all words from a language in a corpus with limited size. And for speed and quality³ considerations, we also don't train word vectors for words which appear less than 5 times in the mono-lingual data. So for those really rare words which are not seen or only seen a few times in the mono-lingual data, we cannot find words similar to them. According to our investigation, most of these really rare words belong to named entities, including number, person names, location names and organization names. With an extra named entity recognizer, we can replace these rare named entities with their type labels instead of similar words. And this will also be leaved as our feature work.

6 Related work

Neural machine translation has a short history of only a few years. [Kalchbrenner and Blunsom, 2013; Cho *et al.*, 2014] first propose to use the encoder-decoder architecture to do sequence to sequence mapping. However, they only use it as an additional feature to evaluate the quality of phrase pairs in traditional machine translation. At the same time, [Sutskever *et al.*, 2014] apply it in end-to-end machine translation. Having considered using only a single vector to represent source sentences with variable lengths is not reasonable, [Bahdanau *et al.*, 2015] propose the attention mechanism to dynamically attend to different source words when generating different target words. [Luong *et al.*, 2015a] propose to use local attention instead of global attention for improved speed and accuracy.

Different to traditional machine translation, NMT model can only employ a small vocabulary due to computational complexity. The rare words problem has attracted a lot of

attention recently. Besides the work by [Luong *et al.*, 2015b] which we compared in our paper, [Jean *et al.*, 2015] propose to directly use large vocabulary with a method based on importance sampling. As pointed out in their paper, their method is complementary and can be used together with replacement methods.

In traditional machine translation, although all vocabulary in the training set can be used for decoding, there are still a lot of out-of-vocabulary words during testing and they hurt the translation performance a lot. Most work [Fung and Cheung, 2004; Marton *et al.*, 2009; Jiang *et al.*, 2007] addressing OOV problem focus on how to translate those OOV correctly during translation. They often resort to additional resources such as comparable data and synonym thesaurus. One notable exception is the work from [Zhang *et al.*, 2012; 2013], which also focuses on the syntactic and semantic role of those OOV and propose to replace OOV with similar words during testing.

7 Conclusion

In this paper, we systematically studied how rare words impact the performance of NMT systems. And we proposed an effective approach of replacing rare words with similar in-vocabulary words. This approach not only enables the translation of rare words, but also reduces the ambiguity of the whole sentence, which is quite important for parameter estimation during training and in-vocabulary words translation during testing. Experiment results on Chinese to English translation tasks demonstrate the power of our methods. Our best replacement method outperforms the baseline by more than 4 BLEU points, which is also much better than the method proposed by previous work.

Acknowledgments

The research work has been partially funded by the Natural Science Foundation of China under Grant No. 61333018 and 91520204 as well, and it is also supported by the Strategic Priority Research Program of the CAS (Grant XDB02070007).

³The word vectors learnt for words with only a few occurrences are not reliable.

References

- [Bahdanau *et al.*, 2015] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. In *ICLR 2015*, 2015.
- [Cho *et al.*, 2014] Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734, Doha, Qatar, October 2014.
- [Fung and Cheung, 2004] Pascale Fung and Percy Cheung. Mining very-non-parallel corpora: Parallel sentence and lexicon extraction via bootstrapping and e. In *EMNLP*, pages 57–63. Citeseer, 2004.
- [Heafield *et al.*, 2013] Kenneth Heafield, Ivan Pouzyrevsky, Jonathan H. Clark, and Philipp Koehn. Scalable modified kneser-ney language model estimation. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 690–696, Sofia, Bulgaria, August 2013.
- [Jean *et al.*, 2015] Sébastien Jean, Kyunghyun Cho, Roland Memisevic, and Yoshua Bengio. On using very large target vocabulary for neural machine translation. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1–10, Beijing, China, July 2015.
- [Jiang *et al.*, 2007] Long Jiang, Ming Zhou, Lee-Feng Chien, and Cheng Niu. Named entity translation with web mining and transliteration. In *IJCAI*, volume 7, pages 1629–1634, 2007.
- [Kalchbrenner and Blunsom, 2013] Nal Kalchbrenner and Phil Blunsom. Recurrent continuous translation models. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1700–1709, Seattle, Washington, USA, October 2013.
- [Koehn *et al.*, 2007] Philipp Koehn, Hieu Hoang, Alexandra Birch, Chris Callison-Burch, Marcello Federico, Nicola Bertoldi, Brooke Cowan, Wade Shen, Christine Moran, Richard Zens, et al. Moses: Open source toolkit for statistical machine translation. In *Proceedings of the 45th annual meeting of the ACL on interactive poster and demonstration sessions*, pages 177–180, 2007.
- [Liang *et al.*, 2006] Percy Liang, Ben Taskar, and Dan Klein. Alignment by agreement. In *Proceedings of the Human Language Technology Conference of the NAACL, Main Conference*, pages 104–111, New York City, USA, June 2006.
- [Luong *et al.*, 2015a] Thang Luong, Hieu Pham, and Christopher D. Manning. Effective approaches to attention-based neural machine translation. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 1412–1421, Lisbon, Portugal, September 2015.
- [Luong *et al.*, 2015b] Thang Luong, Ilya Sutskever, Quoc Le, Oriol Vinyals, and Wojciech Zaremba. Addressing the rare word problem in neural machine translation. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 11–19, Beijing, China, July 2015.
- [Marton *et al.*, 2009] Yuval Marton, Chris Callison-Burch, and Philip Resnik. Improved statistical machine translation using monolingually-derived paraphrases. In *Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing*, pages 381–390, Singapore, August 2009.
- [Mikolov *et al.*, 2013] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119, 2013.
- [Papineni *et al.*, 2002] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA, July 2002.
- [Sutskever *et al.*, 2014] Ilya Sutskever, Oriol Vinyals, and Quoc VV Le. Sequence to sequence learning with neural networks. In *Advances in neural information processing systems*, pages 3104–3112, 2014.
- [Turian *et al.*, 2010] Joseph Turian, Lev Ratinov, and Yoshua Bengio. Word representations: a simple and general method for semi-supervised learning. In *Proceedings of the 48th annual meeting of the association for computational linguistics*, pages 384–394, 2010.
- [Zhang *et al.*, 2012] Jiajun Zhang, Feifei Zhai, and Chengqing Zong. Handling unknown words in statistical machine translation from a new perspective. In *Natural Language Processing and Chinese Computing*, pages 176–187. Springer, 2012.
- [Zhang *et al.*, 2013] Jia-Jun Zhang, Fei-Fei Zhai, and Cheng-Qing Zong. A substitution-translation-restoration framework for handling unknown words in statistical machine translation. *Journal of Computer Science and Technology*, 28(5):907–918, 2013.
- [Zipf, 1949] George Kingsley Zipf. Human behavior and the principle of least effort. 1949.